

Carlos Coronel • Steven Morris

Database Systems

Design, Implementation, & Management

3/19/2024

Database System Design, Implementation And Management

1



National Economics University
Faculty of Mathematical Economics

Database System Design, Implementation And
Management Course – Advanced SQL

- Instructor: Duc Minh Vu (MFE)
- Email: minhvd [at] neu.edu.vn

3/19/2024

Database System Design, Implementation And Management


2

Learning Objectives

- 8-1 Use SQL to create a table manually
- 8-2 Use SQL to create a copy of a table using a subquery
- 8-3 Manipulate the structure of existing tables to add, modify, and remove columns and constraints
- 8-4 Use SQL to do data manipulation (insert, update, and deletions of data)
- 8-5 Use SQL to create database views, including updatable views
- 8-6 Use procedural SQL to create triggers, stored procedures, and procedural SQL functions
- 8-7 Create embedded SQL

8-1 Data Definition Commands

Figure 8.1 Database Model



8-1b Creating the Database

- ❑ If you are using an enterprise RDBMS, you must be authenticated by the RDBMS before you can start creating tables.
- ❑ **Authentication** The process through which a DBMS verifies that only registered users can access the database.
- ❑ To be authenticated, you must log on to the RDBMS using a user ID and a password created by the database administrator.
- ❑ In an enterprise RDBMS, every user ID is associated with a database schema.

8-1c The Database Schema

- ❑ a **schema** is a logical group of database objects—such as tables and indexes—that are related to each other
- ❑ the schema belongs to a single user or application
- ❑ A single database can hold multiple schemas that belong to different users or applications.
- ❑ Schemas in ANSI SQL standards define a command to create a database schema:


```
CREATE SCHEMA AUTHORIZATION {creator};
```

 Therefore, if the creator is JONES, the following command is used:


```
CREATE SCHEMA AUTHORIZATION JONES;
```

Table 8.1 Some Common SQL Data Types

Data Type	Format	Comments
Numeric	NUMERIC(L,D) or NUMBER(L,D)	The declaration NUMERIC(7,2) or NUMBER (7,2) indicates that numbers will be stored with two decimal places and may be up to seven digits long, including the sign and the decimal place (for example, 12.32 or -134.99).
	INTEGER	May be abbreviated as INT. Integers are (whole) counting numbers, so they cannot be used if you want to store numbers that require decimal places.
	SMALLINT	Like INTEGER but limited to integer values up to six digits. If your integer values are relatively small, use SMALLINT instead of INT.
	DECIMAL(L,D)	Like the NUMERIC specification, but the storage length is a <i>minimum</i> specification. That is, greater lengths are acceptable, but smaller ones are not. DECIMAL(9,2), DECIMAL(9), and DECIMAL are all acceptable. In most popular DBMSs, this difference is not implemented so NUMERIC and DECIMAL operate the same.
Character	CHAR(L)	Fixed-length character data for up to 255 characters. If you store strings that are not as long as the CHAR parameter value, the remaining spaces are left unused. Therefore, if you specify CHAR(25), strings such as <i>Smith</i> and <i>Katzenjammer</i> are each stored as 25 characters. However, a U.S. area code is always three digits long, so CHAR(3) would be appropriate if you wanted to store such codes.
	VARCHAR(L) or VARCHAR2(L)	Variable-length character data. The designation VARCHAR(25) or VARCHAR2(25) will let you store up to 25 characters. However, unlike CHAR, VARCHAR will not leave unused spaces. Oracle automatically converts VARCHAR to VARCHAR2.
Date	DATE	Stores dates in the Julian date format.

8-2 Creating Table Structures

8-2a CREATE TABLE command

❑ **CREATE TABLE** A SQL command that creates a table's structures using the characteristics and attributes given.

```
CREATE TABLE tablename (  
    column1          data type          [constraint] [,  
    column2          data type          [constraint] ] [,  
    PRIMARY KEY      (column1          [, column2]) ] [,  
    FOREIGN KEY      (column1          [, column2]) REFERENCES tablename] [,  
    CONSTRAINT       (constraint )];
```


Note

SQL Syntax

Syntax notation for SQL commands used in this book:

CAPITALS	Required SQL command keywords
<i>italics</i>	A parameter provided by the end user (generally required)
{a b ..}	A mandatory parameter; use one option from the list separated by
[.....]	An optional parameter—anything inside square brackets is optional
<i>Tablename</i>	The name of a table
<i>Column</i>	The name of an attribute in a table
<i>data type</i>	A valid data-type definition
<i>constraint</i>	A valid constraint definition
<i>condition</i>	A valid conditional expression (evaluates to true or false)
<i>columnlist</i>	One or more column names or expressions separated by commas
<i>tablelist</i>	One or more table names separated by commas
<i>conditionlist</i>	One or more conditional expressions separated by logical operators
<i>expression</i>	A simple value (such as 76 or Married) or a formula (such as P_PRICE – 10)

3/19/2024

11

```
CREATE TABLE VENDOR (  
  V_CODE      INTEGER          NOT NULL,          UNIQUE,  
  V_NAME      VARCHAR(35)      NOT NULL,  
  V_CONTACT   VARCHAR(25)      NOT NULL,  
  V_AREACODE   CHAR(3)         NOT NULL,  
  V_PHONE     CHAR(8)          NOT NULL,  
  V_STATE     CHAR(2)          NOT NULL,  
  V_ORDER     CHAR(1)          NOT NULL,  
  PRIMARY KEY (V_CODE));
```

3/19/2024

Database System Design, Implementation And Management

12

```
CREATE TABLE PRODUCT (
  P_CODE          VARCHAR(10)          NOT NULL,
  P_DESCRIPT      VARCHAR(35)          NOT NULL,
  P_INDATE        DATE                 NOT NULL,
  P_QOH           SMALLINT             NOT NULL,
  P_MIN           SMALLINT             NOT NULL,
  P_PRICE         NUMBER(8,2)          NOT NULL,
  P_DISCOUNT     NUMBER(5,2)          NOT NULL,
  V_CODE          INTEGER,
  PRIMARY KEY (P_CODE),
  FOREIGN KEY (V_CODE) REFERENCES VENDOR (V_CODE) ON UPDATE
  CASCADE);
```

Note

If you are working with a composite primary key, all of the primary key's attributes are contained within the parentheses and are separated with commas. For example, the LINE table in Figure 8.1 has a primary key that consists of the two attributes INV_NUMBER and LINE_NUMBER. Therefore, you would define the primary key by typing the following:

```
PRIMARY KEY (INV_NUMBER, LINE_NUMBER),
```

The order of the primary key components is important because the indexing starts with the first mentioned attribute, then proceeds with the next attribute, and so on. In this example, the line numbers would be ordered within each of the invoice numbers:

INV_NUMBER	LINE_NUMBER
1001	1
1001	2
1002	1
1003	1
1003	2

The ON UPDATE CASCADE specification ensures that if you make a change in any VENDOR's V_CODE that change is automatically applied to all foreign key references throughout the system to ensure that referential integrity is maintained. (Although the ON UPDATE CASCADE clause is part of the ANSI standard, some RDBMSs, such as Oracle, do not support it. If your RDBMS does not support the clause, delete it from the code shown here.)

Note

Note About Referential Constraint Actions

The support for the referential constraint's actions varies from product to product. For example:

- MySQL requires the InnoDB storage engine to enforce referential integrity.
- MS Access, SQL Server, MySQL, and Oracle support ON DELETE CASCADE.
- MS Access, SQL Server, and MySQL support ON UPDATE CASCADE.
- Oracle does not support ON UPDATE CASCADE.
- Oracle and MySQL support SET NULL.
- MS Access and SQL Server do not support SET NULL.
- Refer to your product manuals for additional information on referential constraints.

While MS Access does not support ON DELETE CASCADE or ON UPDATE CASCADE at the SQL command-line level, it does support them through the relationship window interface. In fact, whenever you try to establish a relationship between two tables in Access, the relationship window interface will automatically pop up.

- The CHECK constraint is used to validate data when an attribute value is entered. The CHECK constraint does precisely what its name suggests: it checks to see that a specified condition exists. Examples of such constraints include the following:
 - The minimum order value must be at least 10.
 - The date must be after April 15, 2018.
 - If the CHECK constraint is met for the specified attribute (i.e., the condition is true), the data is accepted for that attribute. If the condition is found to be false, an error message is generated and the data is not accepted.

```
CREATE TABLE CUSTOMER (  
CUS_CODE          INTEGER          PRIMARY KEY,  
CUS_LNAME         VARCHAR(15)     NOT NULL,  
CUS_FNAME         VARCHAR(15)     NOT NULL,  
CUS_INITIAL       CHAR(1),  
CUS_AREACODE      CHAR(3)         DEFAULT '615' NOT NULL  
CHECK(CUS_AREACODE IN ('615','713','931')),  
CUS_PHONE         CHAR(8)         NOT NULL,  
CUS_BALANCE       DECIMAL(9,2)    DEFAULT 0.00,  
CONSTRAINT CUS_UI1 UNIQUE (CUS_LNAME, CUS_FNAME));
```

In this case, the CUS_AREACODE attribute is assigned a default value of '615'. Therefore, if a new CUSTOMER table row is added and the end user makes no entry for the area code, the '615' value will be recorded. Also, the CHECK condition restricts the values for the customer's area code to 615, 713, and 931; any other values will be rejected.

The DEFAULT value applies only when new rows are added to a table and then only when no value is entered for the customer's area code. (The default value is not used when an existing row in the table is modified.) In contrast, the CHECK condition is validated whether a customer row is added or *modified*. However, while the CHECK condition may include any valid expression, it applies only to the attributes in the table being checked. If you want to check for conditions that include attributes in other tables, you must use triggers, as discussed later in this chapter. Finally, the last line of the CREATE TABLE command sequence creates a unique index constraint (named CUS_UI1) on the customer's last name and first name. The index will prevent the entry of two customers with the same last name and first name. (This index merely illustrates the process. Clearly, it should be possible to have more than one person named John Smith in the CUSTOMER table.)

```
CREATE TABLE INVOICE (  
  INV_NUMBER    INTEGER      PRIMARY KEY,  
  CUS_CODE      INTEGER      NOT NULL REFERENCES CUSTOMER(CUS_CODE),  
  INV_DATE      DATETIME     DEFAULT NOW() NOT NULL,  
  CONSTRAINT INV_CK1 CHECK (INV_DATE > '2022-01-01'));
```

- The CUS_CODE attribute definition contains REFERENCES CUSTOMER(CUS_CODE) to indicate that the CUS_CODE is a foreign key. This is another way to define a foreign key.
- The DEFAULT constraint uses the NOW() function. This function always returns today's date.
- The invoice date (INV_DATE) attribute is automatically given today's date (returned by NOW()) when a new row is added if no value is given for the attribute.
- A CHECK constraint is used to validate that the invoice date is greater than "January 1, 2022."

```
CREATE TABLE LINE (
    INV_NUMBER    INTEGER           NOT NULL,
    LINE_NUMBER   DECIMAL(2,0)      NOT NULL,
    P_CODE        VARCHAR(10)       NOT NULL,
    LINE_UNITS    DECIMAL(9,2)      DEFAULT 0.00    NOT NULL,
    LINE_PRICE    DECIMAL(9,2)      DEFAULT 0.00    NOT NULL,
    PRIMARY KEY (INV_NUMBER, LINE_NUMBER),
    FOREIGN KEY (INV_NUMBER) REFERENCES INVOICE (INV_NUMBER) ON
    DELETE CASCADE,
    FOREIGN KEY (P_CODE) REFERENCES PRODUCT(P_CODE),
    CONSTRAINT LINE_U11 UNIQUE(INV_NUMBER, P_CODE));
```


In the creation of the LINE table, note that a UNIQUE constraint is added to prevent the duplication of an invoice line. A UNIQUE constraint is enforced through the creation of a unique index. Also note that the ON DELETE CASCADE foreign key enforces referential integrity. The use of ON DELETE CASCADE is recommended for weak entities to ensure that the deletion of a row in the strong entity automatically triggers the deletion of the corresponding rows in the dependent weak entity. In that case, the deletion of an INVOICE row will automatically delete all of the LINE rows related to the invoice. In the following section, you will learn more about indexes and how to use SQL commands to create them.

8-2c Creating a Table with a SELECT Statement

❑ SQL provides a way to rapidly create a new table based on selected columns and rows of an existing table using a subquery.

```
CREATE TABLE [IF NOT EXISTS] PART AS
SELECT  P_CODE AS PART_CODE, P_DESCRIPT AS PART_DESCRIPT, P_PRICE
        AS PART_PRICE, V_CODE
FROM    PRODUCT;
```


8-2d SQL Indexes

- ❑ **CREATE INDEX** A SQL command that creates indexes on the basis of a selected attribute or attributes
- ❑ indexes can be used to improve the efficiency of searches and to avoid duplicate column values

- ❑ when you declare a primary key, the DBMS automatically creates a unique index
- ❑ **CREATE [UNIQUE]INDEX indexname ON tablename(column1 [, column2])**

For example, based on the attribute P_INDATE stored in the PRODUCT table, the following command creates an index named P_INDATEX:

```
CREATE INDEX P_INDATEX ON PRODUCT(P_INDATE);
```

8-2d SQL Indexes

- ❑ **DROP INDEX** A SQL command used to delete database objects such as tables, views, indexes, and users.

To delete an index, use the **DROP INDEX** command:

```
DROP INDEX indexname
```

While the ANSI standard for SQL only requires the index name, some DBMSs such as MySQL require specifying the table name as well. For example, if you want to eliminate the PROD_PRICEX index in MySQL, type:

```
DROP INDEX PROD_PRICEX ON PRODUCT;
```

8-3 Altering Table Structures

8-3 Altering Table Structures

❑ **ALTER TABLE** The SQL command used to make changes to table structure. When the command is followed by

a *tablename* and *ADD* or *MODIFY* it adds a column or changes column

ALTER TABLE *tablename*

{ADD | MODIFY | ALTER} [COLUMN] (*columnname datatype* [{ADD | MODIFY |
ALTER} *columnname datatype*]);

ALTER TABLE *tablename*

ADD *constraint* [ADD *constraint*];

ALTER TABLE *tablename*

DROP {PRIMARY KEY | COLUMN *columnname* | CONSTRAINT *constraintname* };

```
ALTER TABLE PRODUCT
    MODIFY V_CODE CHAR(5);

ALTER TABLE PRODUCT
    ADD (P_SALECODE CHAR(1));

ALTER TABLE PART
    ADD
        PRIMARY KEY (PART_CODE);

ALTER TABLE PART
    ADD
        FOREIGN KEY (V_CODE) REFERENCES VENDOR (V_CODE);

ALTER TABLE PART
    ADD
        CHECK (PART_PRICE >= 0);
```

```
ALTER TABLE PART
    ADD
        PRIMARY KEY (PART_CODE),
    ADD
        FOREIGN KEY (V_CODE) REFERENCES VENDOR (V_CODE),
    ADD
        CHECK (PART_PRICE >= 0);

ALTER TABLE LINE
    ADD
        PRIMARY KEY (INV_NUMBER, LINE_NUMBER)
    ADD
        FOREIGN KEY (INV_NUMBER) REFERENCES INVOICE
        (INV_NUMBER)
    ADD
        FOREIGN KEY (P_CODE) REFERENCES PRODUCT (P_CODE);

ALTER TABLE VENDOR
    DROP COLUMN V_ORDER;

DROP TABLE PART;
```

8-4 Data Manipulation Commands

8-4a Adding Table Rows

❑ `INSERT INTO tablename VALUES (value1, value2, ..., valuen)`

`INSERT INTO VENDOR`

`VALUES (21225, 'Bryson, Inc.', 'Smithson', '615', '223-3234', 'TN', 'Y');`

`INSERT INTO VENDOR`

`VALUES (21226, 'Superloo, Inc.', 'Flushing', '904', '215-8995', 'FL', 'N');`

`INSERT INTO PRODUCT(P_CODE, P_DESCRIPT) VALUES ('BRT-345', 'Titanium
drill bit');`

8-4b Inserting Table Rows with a SELECT Subquery

```
INSERT INTO      target_tablename[(target_columnlist)]  
SELECT          source_columnlist  
FROM            source_tablename;
```

```
INSERT INTO PART (PART_CODE, PART_DESCRIPT, PART_PRICE, V_CODE)  
SELECT          P_CODE, P_DESCRIPT, P_PRICE, V_CODE  
FROM PRODUCT;
```

8-4c Saving Table Changes

- ❑ **COMMIT** The SQL command that terminates a transaction by permanently saving data changes to a database.
- ❑ The COMMIT command permanently saves all changes—such as rows added, attributes modified, and rows deleted—made to any table in the database
- ❑ COMMIT [WORK | TRANSACTION]

8-4d Updating Table Rows

❑ **UPDATE** A SQL command that allows attribute values to be changed in one or more

```
ro UPDATE      tablename
SET            columnname = expression [, columnname = expression]
[WHERE        conditionlist];

UPDATE        PRODUCT
SET           P_INDATE = '2022-01-18'
WHERE        P_CODE = '13-Q2/P2';

UPDATE        PRODUCT
SET           P_INDATE = '2022-01-18', P_PRICE = 17.99, P_MIN = 10
WHERE        P_CODE = '13-Q2/P2';
```

8-4e Deleting Table Rows

❑ **DELETE** A SQL command that allows data rows to be deleted from a table.

```
DELETE FROM      tablename
[WHERE           conditionlist];

DELETE FROM      PRODUCT
WHERE            P_CODE = 'BRT-345';
```

8-4f Restoring Table Contents

- ❑ **ROLLBACK** A SQL command that restores the database table contents to the condition that existed after the last COMMIT statement.
- ROLLBACK;
- ❑ If you have not yet used the COMMIT command to store the changes permanently in the database, you can restore the database to its previous condition with the **ROLLBACK** command.
- ❑ COMMIT and ROLLBACK work only with data manipulation commands that add, modify, or delete table rows

COMMIT and ROLLBACK work only with data manipulation commands that add, modify, or delete table rows. For example, assume that you perform these actions:

1. CREATE a table called SALES.
2. INSERT 10 rows in the SALES table.
3. UPDATE two rows in the SALES table.
4. Execute the ROLLBACK command.

Will the SALES table be removed by the ROLLBACK command? No, the ROLLBACK command will undo *only* the results of the INSERT and UPDATE commands. All data definition commands (CREATE TABLE) are automatically committed to the data dictionary and cannot be rolled back. The COMMIT and ROLLBACK commands are examined in greater detail in Chapter 10.

8-5 Virtual Tables: Creating a View

8-5 Virtual Tables: Creating a View

- ❑ **View** A virtual table based on a **SELECT** query that is saved as an object in the database.
- ❑ **base table** The table on which a view is based.
- ❑ **CREATE VIEW** A SQL command that creates a logical, “virtual” table. The view can be treated as a real table.
- ❑ **CREATE VIEW viewname [(column list)] AS SELECT query**



```

CREATE VIEW PROD_STATS AS
SELECT  V_CODE, SUM(P_QOH*P_PRICE) AS TOTCOST, MAX(P_QOH)
        AS MAXQTY, MIN(P_QOH) AS MINQTY, AVG(P_QOH) AS AVGQTY
FROM    PRODUCT
GROUP BY V_CODE;

```

Figure 8.5

Creating a Virtual Table with the Create View Command





8-6 Auto Increment, Identity, and Sequences



8-7 Procedural SQL

8-7 Procedural SQL

- ❑ **persistent storage module (PSM)** A block of code with standard SQL statements and procedural extensions that is stored and executed at the DBMS server.
- ❑ The PSM represents business logic that can be encapsulated, stored, and shared among multiple database users.
- ❑ A PSM lets an administrator assign specific access rights to a stored module to ensure that only authorized users can use it.
- ❑ **Procedural SQL** An extension of the SQL programming language, such as PL/SQL or TSQL, that add procedural programming capabilities, such as variable and logical flow control, to SQL and is designed to run inside the database.

8-7a Stored Procedures

- ❑ **stored procedure (1)** A named collection of procedural and SQL statements
- ❑ They can be used to encapsulate and represent business transactions
- ❑ Stored procedures substantially reduce network traffic and increase performance.
- ❑ Stored procedures help reduce code duplication by means of code isolation and code sharing (creating unique modules that are called by application programs)

```
CREATE PROCEDURE procedure_name ([IN | OUT | INOUT] argument data-type, ...)
BEGIN
    [DECLARE variable data type [DEFAULT initial_value] ]
    [BEGIN]
        Procedural SQL and SQL statements;
    ...
    [END;]
END;
```

- *argument* specifies the parameters that are passed to the stored procedure. A stored procedure could have zero or more arguments or parameters.
- *IN | OUT | INOUT* indicates whether the parameter is for input, output, or both.
- *data-type* is one of the procedural SQL data types used in the RDBMS. The data types normally match those used in the RDBMS table creation statement.
- The body of the stored procedure must be inside a BEGIN...END; block. Within that block can be one or more DECLARE statements and other blocks of executable code.
- Variables can be declared in a DECLARE statement. You must specify the variable name, its data type, and (optionally) an initial value.

Figure 8.9 Block Program with Insert

```
1 delimiter $$
2 DROP PROCEDURE IF EXISTS TEMP_PROC$$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     INSERT INTO VENDOR VALUES (25678, 'Microsoft Corp.', 'Bill Gates', '765', '546-8484', 'WA', 'N');
6 END;
7 $$
8 delimiter ;
9
10 call temp_proc()
```

Figure 8.10 Output to the Screen

```
1 delimiter $$
2 DROP PROCEDURE IF EXISTS TEMP_PROC$$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     INSERT INTO VENDOR VALUES (25678, 'Microsoft Corp.', 'Bill Gates', '765', '546-8484', 'WA', 'N');
6     SELECT 'New Vendor Added!' as 'Result';
7 END;
8 $$
9 delimiter ;
10
11 call temp_proc()
12
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: 1A

Result
New Vendor Added!

8-7b Working with Variables

Figure 8.11 Variable Manipulation

```
1 DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     DECLARE MYNUM1 NUMERIC(1) DEFAULT 5;
6     DECLARE MYNUM2 NUMERIC(1);
7     DECLARE TOTAL INT;
8     DECLARE MESSAGE VARCHAR(15) DEFAULT 'The total is ';
9
10    SET MYNUM2 = 6;
11    SET TOTAL = MYNUM1 * MYNUM2;
12    SELECT CONCAT(MESSAGE, TOTAL) AS 'Result';
13 END;
14 $$
15 delimiter ;
16
17 call temp_proc();
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: IA

Result
The total is 30

Figure 8.12 Conditional Execution

```
1 • DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 • CREATE PROCEDURE TEMP_PROC()
4   BEGIN
5     DECLARE MYNUM1 INT DEFAULT 5;
6     DECLARE MYNUM2 INT DEFAULT 60;
7     DECLARE TOTAL INT;
8     DECLARE MESSAGE VARCHAR(50);
9
10    SET TOTAL = MYNUM1 * MYNUM2;
11    IF TOTAL < 100 THEN
12      SET MESSAGE = 'The total is small: ' ;
13    ELSE
14      SET MESSAGE = 'The total is large: ' ;
15    END IF;
16    SELECT CONCAT(MESSAGE, TOTAL) AS 'Result';
17  END;
18 $$
19 delimiter ;
20
21 • call temp_proc();
```

Result
The total is large: 300

Figure 8.13 ELSEIF for Complex IF Structure

```
1 • DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 • CREATE PROCEDURE TEMP_PROC()
4   BEGIN
5     DECLARE MYNUM1 INT DEFAULT 50;
6     DECLARE MYNUM2 INT DEFAULT 10;
7     DECLARE TOTAL INT;
8     DECLARE MESSAGE VARCHAR(50);
9     SET TOTAL = MYNUM1 * MYNUM2;
10    IF TOTAL < 100 THEN
11      SET MESSAGE = 'The total is small: ' ;
12    ELSEIF TOTAL >= 100 AND TOTAL <= 500 THEN
13      SET MESSAGE = 'The total is medium: ' ;
14    ELSEIF TOTAL BETWEEN 500 AND 1000 THEN
15      SET MESSAGE = 'The total is big: ' ;
16    ELSE
17      SET MESSAGE = 'The total is huge: ' ;
18    END IF;
19    SELECT CONCAT(MESSAGE, TOTAL) AS 'Result';
20  END;
21 $$
22 delimiter ;
23 • call temp_proc();
```

Result
The total is medium: 500

Figure 8.14 A Simple Loop

```

1 DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     DECLARE MYNUM INT DEFAULT 1;
6     DECLARE RESULT VARCHAR(100) DEFAULT MYNUM;
7     COUNTER: LOOP
8         SET MYNUM = MYNUM + 1;
9         IF MYNUM > 20 THEN
10             LEAVE COUNTER;
11         END IF;
12         SET RESULT = CONCAT(RESULT, ', ', MYNUM);
13     END LOOP;
14     SELECT RESULT;
15 END;
16 $$
17 delimiter ;
18 call temp_proc();

```

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
RESULT			
1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20			

Figure 8.15 Implementation of a WHILE Loop

```

1 DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     DECLARE MYNUM INT DEFAULT 1;
6     DECLARE RESULT VARCHAR(100) DEFAULT MYNUM;
7     WHILE MYNUM < 20 DO
8         SET MYNUM = MYNUM + 1;
9         SET RESULT = CONCAT(RESULT, ', ', MYNUM);
10     END WHILE;
11     SELECT RESULT;
12 END;
13 $$
14 delimiter ;
15 call temp_proc();

```

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
RESULT			
1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20			

Figure 8.17 Temporary Table for Outputs

```
1 DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     DECLARE MYNUM INT DEFAULT 1;
6     CREATE TEMPORARY TABLE MYTAB(RESULT VARCHAR(100));
7     WHILE MYNUM <= 5 DO
8         INSERT INTO MYTAB VALUES (CONCAT('The current value is ', MYNUM));
9         SET MYNUM = MYNUM + 1;
10    END WHILE;
11    SELECT * FROM MYTAB;
12    DROP TEMPORARY TABLE MYTAB;
13 END;
14 $$
15 delimiter ;
16 call temp_proc();
```

RESULT
The current value is 1
The current value is 2
The current value is 3
The current value is 4
The current value is 5

Figure 8.19 An Explicit Cursor Example

```
1 DROP PROCEDURE IF EXISTS temp_proc;
2 delimiter $$
3 CREATE PROCEDURE TEMP_PROC()
4 BEGIN
5     DECLARE M_P_CODE VARCHAR(15);
6     DECLARE M_P_DESCRPT VARCHAR(50);
7     DECLARE M_TOT INT DEFAULT 0;
8     DECLARE NOTFOUND INT DEFAULT 0;
9     DECLARE PROD_CURSOR CURSOR FOR SELECT P_CODE, P_DESCRPT
10    FROM PRODUCT
11    WHERE P_QOH > (SELECT AVG(P_QOH) FROM PRODUCT);
12    DECLARE CONTINUE HANDLER FOR NOT FOUND SET NOTFOUND = 1;
13    DROP TEMPORARY TABLE IF EXISTS MYTAB;
14    CREATE TEMPORARY TABLE MYTAB(RESULT VARCHAR(100));
15    OPEN PROD_CURSOR;
16    INSERT INTO MYTAB VALUES ('-----');
17    GET_PRODUCTS: LOOP
18        FETCH PROD_CURSOR INTO M_P_CODE, M_P_DESCRPT;
19        IF NOTFOUND = 1 THEN
20            LEAVE GET_PRODUCTS;
21        END IF;
22        INSERT INTO MYTAB VALUES (CONCAT(M_P_CODE, ' -> ', M_P_DESCRPT));
23        SET M_TOT = M_TOT + 1;
24    END LOOP;
25    CLOSE PROD_CURSOR;
26    INSERT INTO MYTAB VALUES ('-----');
27    INSERT INTO MYTAB VALUES (CONCAT('TOTAL PRODUCTS PROCESSED: ', M_TOT));
28    INSERT INTO MYTAB VALUES ('--- END OF REPORT ---');
29    SELECT RESULT AS "PRODUCTS WITH P_QOH > AVG(P_QOH)" FROM MYTAB;
30    DROP TEMPORARY TABLE MYTAB;
31 END;
32 $$
33 delimiter ;
34 call temp_proc();
```

PRODUCTS WITH P_QOH > AVG(P_QOH)

PIC200RT -> PIC 000, 2.5m, 8ft
SH-30277 -> 1.25m metal screw, 25
SH-23118 -> 2.5m yd screw, 50

TOTAL PRODUCTS PROCESSED: 3

--- END OF REPORT ---

Figure 8.21 Second Version of the PRC_PROD_DISCOUNT Stored Procedure

```

1 DROP PROCEDURE IF EXISTS prc_prod_discount;
2 delimiter $$
3 CREATE PROCEDURE PRC_PROD_DISCOUNT(IN WPI DECIMAL(2,2))
4 BEGIN
5     IF WPI <= 0 OR WPI >= 1 THEN -- validate WPI parameter
6         SELECT 'Error: Value must be greater than 0 and less than 1' AS RESULT;
7     ELSE
8         UPDATE PRODUCT
9         SET P_DISCOUNT = P_DISCOUNT + WPI
10        WHERE P_QOH >= P_MIN * 2;
11        SELECT '*** Update finished.';
12    END IF;
13 END;
14 $$
15
16
17 START TRANSACTION;
18 CALL PRC_PROD_DISCOUNT(.33);
19 SELECT P_CODE, P_DESCRIPTION, P_QOH, P_MIN, P_DISCOUNT FROM PRODUCT;
20 ROLLBACK;

```

P_CODE	P_DESCRIPTION	P_QOH	P_MIN	P_DISCOUNT
11QER-21	Power painter, 15 psi., 3-nozzle	8	5	0.00
13-Q2/P2	7.25-in. pwr. saw blade	32	15	0.38
14-Q1/L3	9.00-in. pwr. saw blade	18	12	0.00
15-46-QQ2	Hrd. cloth, 1/4-in., 2x50	15	8	0.00
1558-QW1	Hrd. cloth, 1/2-in., 3x50	23	5	0.33
2232-QTY	880 g/saw, 12-in. blade	8	5	0.05
2232-QWE	880 g/saw, 8-in. blade	6	5	0.05
2238/QPD	880 cordless drill, 1/2-in.	12	5	0.38
23109-HB	Claw hammer	23	10	0.43
23114-AA	Sledge hammer, 12 lb.	8	5	0.05
54778-ZT	Rat-tail file, 1/8-in. fine	43	20	0.33
89-WRE-Q	Hcut chain saw, 16 in.	11	5	0.38
PVC23CR7	PVC pipe, 3.5-in., 8-ft.	188	75	0.33
SM-18277	1.25-in. metal screw, 25	172	75	0.33
SW-23116	2.5-in. wd. screw, 50	237	100	0.33
VR3/TT3	Steel matting, 48x116", 5" m...	18	5	0.43

8-7g Triggers

- ❑ **Trigger** A procedural SQL code that is automatically invoked by the relational database management system when a data manipulation event occurs.
- ❑ A trigger is invoked before or after a data row is inserted, updated, or deleted.
- ❑ A trigger is associated with a database table.
- ❑ Each database table may have one or more triggers.
- ❑ A trigger is executed as part of the transaction that triggered it

8-7g Triggers (cont.)

- ❑ For example, if a product's quantity on hand is updated when the product is sold, the system should automatically check whether the quantity on hand falls below its minimum allowable quantity.
- ❑ To demonstrate that process, use the `PRODUCT` table in Figure 8.2.2. Note the use of the minimum quantity (`P_MIN`) and product reorder flag (`P_REORDER`) columns.
- ❑ The `P_MIN` indicates the minimum quantity on hand allowable before restocking an item.
- ❑ The `P_REORDER` column is a numeric field that indicates whether the product needs to be reordered (1 – Yes, – 5 No).
- ❑ The initial `P_REORDER` values are set to 0 (No) to serve as the basis for the initial

```
CREATE TRIGGER trigger_name
[BEFORE / AFTER] [DELETE / INSERT / UPDATE] ON table_name
FOR EACH ROW
[DECLARE]
[variable_namedata type [=initial_value] ]
BEGIN
Procedural SQL instructions;
...
END;
```

Review Questions


1. What type of integrity is enforced when a primary key is declared?
2. Explain why it might be more appropriate to declare an attribute that contains only digits as a character data type instead of a numeric data type.
3. What is the difference between a column constraint and a table constraint?
4. What are “referential constraint actions”?
5. What is the purpose of a CHECK constraint?
6. Explain when an ALTER TABLE command might be needed.
7. What is the difference between an INSERT command and an UPDATE command?
8. What is the difference between using a subquery with a CREATE TABLE command and using a subquery with an INSERT command?
9. What is a sequence? Write its syntax.
10. What is a trigger, and what is its purpose? Give an example.
11. What is a stored procedure, and why is it particularly useful? Give an example.

Problems

The Ch08_ConstructCo database stores data for a consulting company that tracks all charges to projects. The charges are based on the hours each employee works on each project. The structure and contents of the Ch08_ConstructCo database are shown in Figure P8.1.

Figure P8.1 The Ch08_ConstructCo Database

Relational diagram



Database name: Ch08_ConstructCo

Table name: EMPLOYEE

EMP_NUM	EMP_INITIAL	EMP_FNAME	EMP_LNAME	EMP_HIREDATE	JOB_CODE	EMP_YEARS
101	News	John	G	08-Nov-04 502		17
102	Senior	David	H	12-Jul-91 501		28
103	Artbrough	June	E	01-Dec-00 500		21
104	Ramores	Anne	K	15-Nov-91 501		30
105	Johnson	Alice	K	01-Feb-97 502		25
106	Smithfield	William	K	22-Jan-08 500		16
107	McGee	John	D	15-Jun-99 501		24
108	Washington	Ralph	B	23-Aug-95 501		26
109	Smith	Larry	W	18-Jul-01 501		20
110	Olenko	Gerald	A	11-Dec-99 505		22
111	Walsh	Geoff	B	04-Apr-95 506		27
112	Smithson	Darlene	M	23-Oct-96 507		23
113	Jamthoood	Debbie	K	15-Nov-00 508		21
114	Jones	Annalise	L	20-Aug-97 508		24
115	Benavigi	Travis	B	05-Mar-01 510		21
116	Williamson	Gerald	L	15-Jun-99 501		24
117	Forrester	James	J	04-Jan-99 510		21
118	Forrester	James	J	04-Jan-99 510		11

Table name: ASSIGNMENT

ASSIGN_NUM	ASSIGN_CODE	PROJ_NUM	EMP_NUM	ASSIGN_DATE	ASSIGN_HOUR	ASSIGN_AMOUNT	ASSIGN_COST
1001	23-Mar-22	21	117	2000	34.15	4.2	145.18
1002	23-Mar-22	18	117	2000	34.15	2	89.5
1003	23-Mar-22	18	117	2000	34.15	2	89.5
1004	23-Mar-22	25	108	2000	34.15	3.2	222.204
1005	23-Mar-22	27	104	2000	34.15	4.2	406.104
1006	23-Mar-22	27	113	2000	34.15	3.8	353.8
1007	23-Mar-22	27	113	2000	34.15	3.8	353.8
1008	23-Mar-22	27	113	2000	34.15	3.8	353.8
1009	23-Mar-22	27	113	2000	34.15	3.8	353.8
1010	23-Mar-22	27	113	2000	34.15	3.8	353.8
1011	23-Mar-22	27	113	2000	34.15	3.8	353.8
1012	23-Mar-22	27	113	2000	34.15	3.8	353.8
1013	23-Mar-22	27	113	2000	34.15	3.8	353.8
1014	23-Mar-22	27	113	2000	34.15	3.8	353.8
1015	23-Mar-22	27	113	2000	34.15	3.8	353.8
1016	23-Mar-22	27	113	2000	34.15	3.8	353.8
1017	23-Mar-22	27	113	2000	34.15	3.8	353.8
1018	23-Mar-22	27	113	2000	34.15	3.8	353.8
1019	23-Mar-22	27	113	2000	34.15	3.8	353.8
1020	23-Mar-22	27	113	2000	34.15	3.8	353.8

Table name: PROJECT

PROJ_NUM	PROJ_NAME	PROJ_VALUE	PROJ_BALANCE	EMP_NUM
1	Engineer	1555000	1555000	110
2	Engineer	1555000	1555000	110
3	Engineer	1555000	1555000	110
4	Engineer	1555000	1555000	110
5	Engineer	1555000	1555000	110
6	Engineer	1555000	1555000	110
7	Engineer	1555000	1555000	110
8	Engineer	1555000	1555000	110
9	Engineer	1555000	1555000	110
10	Engineer	1555000	1555000	110
11	Engineer	1555000	1555000	110
12	Engineer	1555000	1555000	110
13	Engineer	1555000	1555000	110
14	Engineer	1555000	1555000	110
15	Engineer	1555000	1555000	110
16	Engineer	1555000	1555000	110
17	Engineer	1555000	1555000	110
18	Engineer	1555000	1555000	110
19	Engineer	1555000	1555000	110
20	Engineer	1555000	1555000	110
21	Engineer	1555000	1555000	110
22	Engineer	1555000	1555000	110
23	Engineer	1555000	1555000	110
24	Engineer	1555000	1555000	110
25	Engineer	1555000	1555000	110

Note that the ASSIGNMENT table in Figure P8.1 stores the JOB_CHG_HOUR values as an attribute (ASSIGN_CHG_HR) to maintain historical accuracy of the data. The JOB_CHG_HOUR values are likely to change over time. In fact, a JOB_CHG_HOUR change will be reflected in the ASSIGNMENT table. Naturally, the employee primary job assignment might also change, so the ASSIGN_JOB is also stored. Because those attributes are required to maintain the historical accuracy of the data, they are *not* redundant.

Given the structure and contents of the Ch08_ConstructCo database shown in Figure P8.1, use SQL commands to answer Problems 1–15.

1. Write the SQL code that will create only the table structure for a table named EMP_1. This table will be a subset of the EMPLOYEE table. The basic EMP_1 table structure is summarized in the following table. Use EMP_NUM as the primary key. Note that the JOB_CODE is the FK to JOB so be certain to enforce referential integrity. Your code should also prevent null entries in EMP_LNAME and EMP_FNAME.

Table P8.1 ConstructCo database structure

Attribute (Field) Name	Data Declaration
EMP_NUM	VARCHAR(3)
EMP_LNAME	VARCHAR(15)
EMP_FNAME	VARCHAR(15)
EMP_INITIAL	CHAR(1)
EMP_HIREDATE	DATE
JOB_CODE	VARCHAR(3)

- Having created the table structure in Problem 1, write the SQL code to enter the first two rows for the table shown in Figure P8.2. Each row should be inserted individually, without using a subquery. Insert the rows in the order that they are listed in the figure.

Figure P8.2 The Contents of the EMP_1 Table

EMP_NUM	EMP_LNAME	EMP_FNAME	EMP_INITIAL	EMP_HIREDATE	JOB_CODE
101	News	John	G	08-Nov-04	502
102	Senior	David	H	12-Jul-93	501
103	Arbough	June	E	01-Dec-00	500
104	Ramoras	Anne	K	15-Nov-91	501
105	Johnson	Alice	K	01-Feb-97	502
107	Alonzo	Maria	D	10-Oct-97	501
108	Washington	Ralph	B	22-Aug-95	501
109	Smith	Larry	W	18-Jul-01	501

- Using the EMPLOYEE table that already exists, use a subquery to insert the remaining rows from the EMPLOYEE table into the EMP_1 table. Remember, your subquery should only retrieve the columns needed for the EMP_1 table and only the employees shown in the figure.
- Write the SQL code that will save the changes made to the EMP_1 table (if supported by your DBMS).
- Write the SQL code to change the job code to 501 for the person whose employee number (EMP_NUM) is 107.
- Write the SQL code to delete the row for William Smithfield, who was hired on June 22, 2008, and whose job code is 500. (*Hint:* Use logical operators to include all of the information given in this problem. Remember, if you are using MySQL, you will have to first disable “safe mode.”)
- Write the SQL code to create a copy of EMP_1, including all of its data, and naming the copy EMP_2.

- Using the EMP_2 table, write the SQL code that will add the attributes EMP_PCT and PROJ_NUM to EMP_2. The EMP_PCT is the bonus percentage to be paid to each employee. The new attribute characteristics are:

EMP_PCT	DECIMAL(4,2)
PROJ_NUM	CHAR(3)

- Using the EMP_2 table, write the SQL code to change the EMP_PCT value to 3.85 for the person whose employee number (EMP_NUM) is 103.
- Using the EMP_2 table, write a single SQL command to change the EMP_PCT value to 5.00 for the people with employee numbers 101, 105, and 107.
- Using the EMP_2 table, write a single SQL command to change the EMP_PCT value to 10.00 for all employees who do not currently have a value for EMP_PCT.
- Using the EMP_2 table, write the SQL command to add .15 to the EMP_PCT of the employee whose name is Maria D. Alonzo. (Use the employee name in your command to determine the correct employee.)
- Using a single command sequence with the EMP_2 table, write the SQL code that will change the project number (PROJ_NUM) to 18 for all employees whose job classification (JOB_CODE) is 500.
- Using a single command sequence with the EMP_2 table, write the SQL code that will change the project number (PROJ_NUM) to 25 for all employees whose job classification (JOB_CODE) is 502 or higher.
- Write the SQL code that will change the PROJ_NUM to 14 for employees who were hired before January 1, 1998, and whose job code is at least 501. When you finish Problems 7–15, the EMP_2 table will contain the data shown in Figure P8.15.

Figure P8.15 The EMP_2 Table Contents

EMP_NUM	EMP_LNAME	EMP_FNAME	EMP_INITIAL	EMP_HIREDATE	JOB_CODE	EMP_PCT	PROJ_NUM
101	News	John	G	08-Nov-00	502	5	25
102	Senior	David	H	12-Jul-89	501	10	14
103	Arbough	June	E	01-Dec-96	500	3	85
104	Ranoras	Anne	K	15-Nov-87	501	10	14
105	Johnson	Alice	K	01-Feb-93	502	5	14
107	Alonzo	Maria	D	10-Oct-93	501	5	15
108	Washington	Ralph	B	22-Aug-91	501	10	14
109	Smith	Larry	W	18-Jul-97	501	10	

Use the database tables in Figure P8.16 as the basis for Problems 16–30.

Figure P8.16 Ch08_SimpleCo Database Tables

Table name: CUSTOMER				Table name: INVOICE			
CUST_NUM	CUST_LNAME	CUST_FNAME	CUST_BALANCE	INV_NUM	CUST_CODE	INV_DATE	INV_AMOUNT
1000	Smith	Jeanne	1050.11	8000	1000	23-Mar-22	235.89
1001	Ortega	Juan	840.92	8001	1001	23-Mar-22	312.82
				8002	1001	30-Mar-22	528.10
				8003	1000	12-Apr-22	194.78
				8004	1000	23-Apr-22	619.44

16. Create the CUSTOMER table structure illustrated in Figure P8.16. The customer number should store integer values. The name attributes should support variable length character data up to 30 characters each. The customer balance should support up to six digits on the left of the decimal place and two digits to the right of the decimal place.
17. Create the INVOICE table structure illustrated in Figure P8.16. The invoice number should store integer values. The invoice date should store date values. The invoice amount should support up to 8 digits to the left of the decimal place and two digits to the right of the decimal place.
18. Write the set of SQL commands necessary to insert the data into the CUSTOMER table you created in Problem 16, as illustrated in Figure P8.16.
19. Write the set of SQL commands necessary to insert the data into the INVOICE table you created in Problem 17, as illustrated in Figure P8.16.
20. Enable auto increment to generate values for customer numbers. The values should start with 2000.
21. Enable auto increment to generate values for invoice numbers. The values should start with the value 9000.
22. Insert the following customer into the CUSTOMER table, using the auto increment from Problem 20 to generate the customer number automatically:
'Powers', 'Ruth', 500
23. Modify the CUSTOMER table to include the customer's date of birth (CUST_DOB), which should store date data.
24. Modify customer 1000 to indicate the date of birth as March 15, 1989.
25. Modify customer 1001 to indicate the date of birth as December 22, 1988.

26. Create a trigger named **trg_updatecustbalance** to update the CUST_BALANCE in the CUSTOMER table when a new invoice record is entered. (Assume that the sale is a credit sale.) Whatever value appears in the INV_AMOUNT column of the new invoice should be added to the customer's balance. Test the trigger using the following new INVOICE record, which would add 225.40 to the balance of customer 1001:

8005, 1001, '2022-04-27', 225.40

27. Write a procedure named **prc_cust_add** to add a new customer to the CUSTOMER table. Use the following values in the new record:

1002, 'Raulthor', 'Peter', 0.00

(You should execute the procedure and verify that the new customer was added to ensure your code is correct.)

28. Write a procedure named **prc_invoice_add** to add a new invoice record to the INVOICE table. Use the following values in the new record:

8006, 1000, '2022-04-30', 301.72

(You should execute the procedure and verify that the new invoice was added to ensure your code is correct.)

29. Write a trigger to update the customer balance when an invoice is deleted. Name the trigger **trg_updatecustbalance2**.

30. Write a procedure to delete an invoice, giving the invoice number as a parameter. Name the procedure **prc_inv_delete**. Test the procedure by deleting invoices 8005 and 8006.

Use the Ch08_SaleCo2 database shown in Figure P8.31 to work Problems 31–34.

Figure P8.31 Ch08_SaleCo2 Database Tables

Database name: Ch08_SaleCo

Table name: INVOICE

Table name: CUSTOMER

Table name: VENDOR

Table name: LINE

Table name: PRODUCT

35. Modify the MODEL table to add the attribute and insert the values shown in the following table.

Table P8.35 Attribute and Values for Problem 35

Attribute Name	Attribute Description	Attribute Type	Attribute Values
MOD_WAIT_CHG	Waiting charge per hour for each model	Numeric	\$100 for C-90A \$50 for PA23-250 \$75 for PA31-350

36. Write the queries to update the MOD_WAIT_CHG attribute values based on Problem 35.
37. Modify the CHARTER table to add the attributes shown in the following table.

Table P8.37 Attributes and Values for Problem 37

Attribute Name	Attribute Description	Attribute Type
CHAR_WAIT_CHG	Waiting charge for each model (copied from the MODEL table)	Numeric
CHAR_FLT_CHG_HR	Flight charge per mile for each model (copied from the MODEL table using the MOD_CHG_MILE attribute)	Numeric
CHAR_FLT_CHG	Flight charge (calculated by $\text{CHAR_HOURS_FLOWN} \times \text{CHAR_FLT_CHG_HR}$)	Numeric
CHAR_TAX_CHG	$\text{CHAR_FLT_CHG} \times \text{tax rate (8\%)}$	Numeric
CHAR_TOT_CHG	$\text{CHAR_FLT_CHG} + \text{CHAR_TAX_CHG}$	Numeric
CHAR_PMT	Amount paid by customer	Numeric
CHAR_BALANCE	Balance remaining after payment	Numeric

38. Write the sequence of commands required to update the CHAR_WAIT_CHG attribute values in the CHARTER table.
(Hint: This could be done with an update, an updatable view, or a stored procedure.)
39. Write the sequence of commands required to update the CHAR_FLT_CHG_HR attribute values in the CHARTER table.
40. Write the command required to update the CHAR_FLT_CHG attribute values in the CHARTER table.
41. Write the command required to update the CHAR_TAX_CHG attribute values in the CHARTER table.
42. Write the command required to update the CHAR_TOT_CHG attribute values in the CHARTER table.
43. Modify the PILOT table to add the attribute shown in the following table.

Table P8.43 Attribute to Add in Problem 43

Attribute Name	Attribute Description	Attribute Type
PILOT_PIC_HRS	Pilot in command (PIC) hours; updated by adding the CHARTER table's CHAR_HOURS_FLOWN to the PILOT_PIC_HRS when the CREW table shows the CREW_JOB to be Pilot	Numeric

44. Create a trigger named **trg_char_hours** that automatically updates the AIRCRAFT table when a new CHARTER row is added. Use the CHARTER table's CHAR_HOURS_FLOWN to update the AIRCRAFT table's AC_TTAF, AC_TTEL, and AC_TTER values.

Chapter 8: Advanced SQL 425

45. Create a trigger named **trg_pic_hours** that automatically updates the **PILOT** table when a new **CREW** row is added and the **CREW** table uses a **Pilot CREW_JOB** entry. Use the **CHARTER** table's **CHAR_HOURS_FLOWN** to update the **PILOT** table's **PIL_PIC_HRS** only when the **CREW** table uses a **Pilot CREW_JOB** entry.
46. Create a trigger named **trg_cust_balance** that automatically updates the **CUSTOMER** table's **CUS_BALANCE** when a new **CHARTER** row is added. Use the **CHARTER** table's **CHAR_TOT_CHG** as the update source. (Assume that all charter charges are charged to the customer balance.)